


# Balanced Binary Search Trees

An Idea To Avoid Worst Case Binary Search Tree (Degeneration Into List)

# Limitations of Ordinary BST

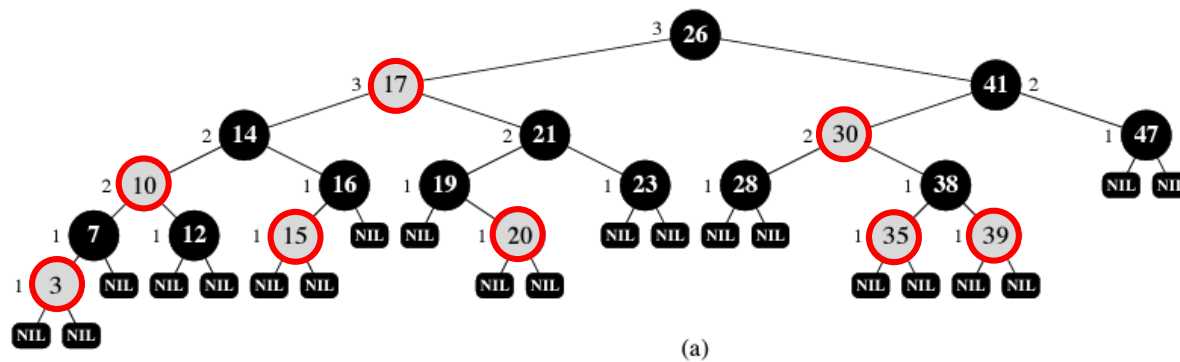
- $O(h)$  time complexities for all core operations (search/insert/delete/...)
- What is the resulting *ordinary* BST when values are inserted in the following order?  
10, 20, 30, 40, 50, 60, 70
- What is the *optimal* BST for the above input sequence?
  - Optimal in the sense of smallest height  $h$ .




# Red-Black Tree

- Special kind of BST with additional node property (color=red/black only) and following additional requirements so that the worst case tree height is still guaranteed to be  $O(\lg n)$ , not  $O(n)$ :
  1. Every node is either red or black.
  2. The root is black.
  3. Every leaf (NIL) is black. A value-bearing node is NOT considered a leaf.
  4. If a node is red, then both its children are black.
    - No two consecutive reds along any simple downward path.
  5. For each node, all simple paths from the node to descendant leaves contain the same number of black nodes.
    - $bh(x)$ : Black-height of node  $x$ , denoting # black nodes on any simple path from  $x$ , not including  $x$  itself

# Red-Black Tree Example (CLRS pp.310)



Black-height is denoted on each node's side.



NIL leaves are usually omitted, but remember that they still contribute to black-heights!


# How To Distinguish Non-Red-Black Tree

- Requirements 1, 2, 3 are trivial:
  1. Every node is either red or black.
  2. Root is black.
  3. Every leaf (NIL) is black.
- Requirement 4 is not too hard:
  4. If a node is red, then both its children are black.
- Requirement 5 is probably the hardest to check:
  5. For each node, all simple paths from the node to descendant leaves contain the same number of black nodes
    - Technique: Calculate/write black-height from bottom-up!


# Non-Red-Black Tree Examples

26

Non-black root (violating Req. 2)



26's left black-height is not equal to its right black-height  
(violating Req. 5)



Red having red child (violating Req. 4)




Not even a BST (violating the whole premise)

# Balanced Nature of Red-Black Trees

- Lemma (CLRS pp.309): A red-black tree with  $n$  internal nodes has height at most  $2 \lg(n + 1)$ .
  - Sub-lemma: The subtree rooted at any node  $x$  has at least  $2^{bh(x)} - 1$  internal nodes.
    - Proof by induction: Read/understand carefully.
    - There's an intuitive understanding to this lemma, though.
  - Another observation: The black-height of the root must be at least  $h/2$ .
    - Because of requirement 4, at least half the nodes on any simple path from the root to a leaf must be black.
  - Thus, applying the above sub-lemma on the root, we get:
    - $n \geq 2^{bh(r)} - 1 \geq 2^{h/2} - 1$
    - Moving 1 and taking logarithms, we get  $h \leq 2 \lg(n + 1)$ .

# Intuition On Red-Black Tree's Height Bound

- Worst case left/right height difference (skewness) could be at most twice!



- You can't add any child to 2 or 4, without first adding other nodes on other parts of the red-black tree, thus limiting any more deviation!



# Inserting in a Red-Black Tree

Insert As Done On An Ordinary BST, Then Fix To Recover Red-Black Properties

# Lesson Objectives


- Identify the resulting red-black tree after inserting an arbitrary new value into a given red-black tree.
- Given an incomplete Red-Black tree insertion code, fill in the blanks for correct Red-Black tree insertion operations.

# Inserting New Value, Retaining R-B Properties

- Recall inserting 10, 20, 30, 40, 50, 60, 70 in sequence to an ordinary BST
- What if the BST needs to be R-B tree all the time?
  - Find the insertion spot, treating the given R-B tree as an ordinary BST
  - Determine the color of the inserted node
    - So that the chosen color doesn't violate the R-B properties
  - If neither choice is possible,
    - Fix the tree structure!

# Trivial Insertions

- Inserting into an empty R-B tree:
  - The inserted value occupies the root node, which must be colored black.
- If insertion location is a black node's child:
  - The inserted node can be simply colored red, without violating any R-B properties (no consecutive reds, same black heights everywhere).




# Nontrivial Insertions

- If insertion location is a red node's child:
  - The newly inserted node can't be black (matching-left-and-right-black-heights rule broken).
  - The newly inserted node can't be red either (no-consecutive-reds rule broken).
  - Then how?
    - Keep one rule that's harder to enforce (same black heights), and fix the other rule that's broken (no consecutive reds).
    - That is, insert as a red node, breaking the no-consecutive-reds rule, and fix it along the path to root.




# More Nontrivial Insertions


- Insert 40:
  - Rotate along 20 & recolor doesn't work, because 40's uncle is red.
  - Instead of rotating, we can simply flip colors of 40's parent, uncle & grandparent, and yet still meeting the same-black-heights rule!
  - Then the grandparent and the great grandparent need to be checked for consecutive reds, and this process repeats (bubbling up).
  - If there's no great grandparent, then the grandparent is root, which is now red, but can be simply repainted black, without violating the same-black-heights rule.




- Insert 50: Just rotate-and-recolor (50's uncle is not red, but black).



- Inserting 60, 70, 80, ... : Same pattern, but may bubble up.




- Inserting 75: Just to show a different situation (Case 2 in CLRS)—We can easily transform it to a well-known case (Case 3 in CLRS).






## BST Rotation Example (CLRS Fig. 13.3)




# Rotating Red-Black Tree


- When to rotate:
  - For the red violating node ( $z$ ), its parent is red, its uncle is black.
  - Its grandparent must be always black.
- CLRS Fig. 13.6 ( $\delta$  is a subtree whose root is black)



Case 2




Case 3




- After this process, there's no more consecutive-reds-violation!

# Red-Black Tree Rotation (Fix-Up) Example (CLRS Fig. 13.4)



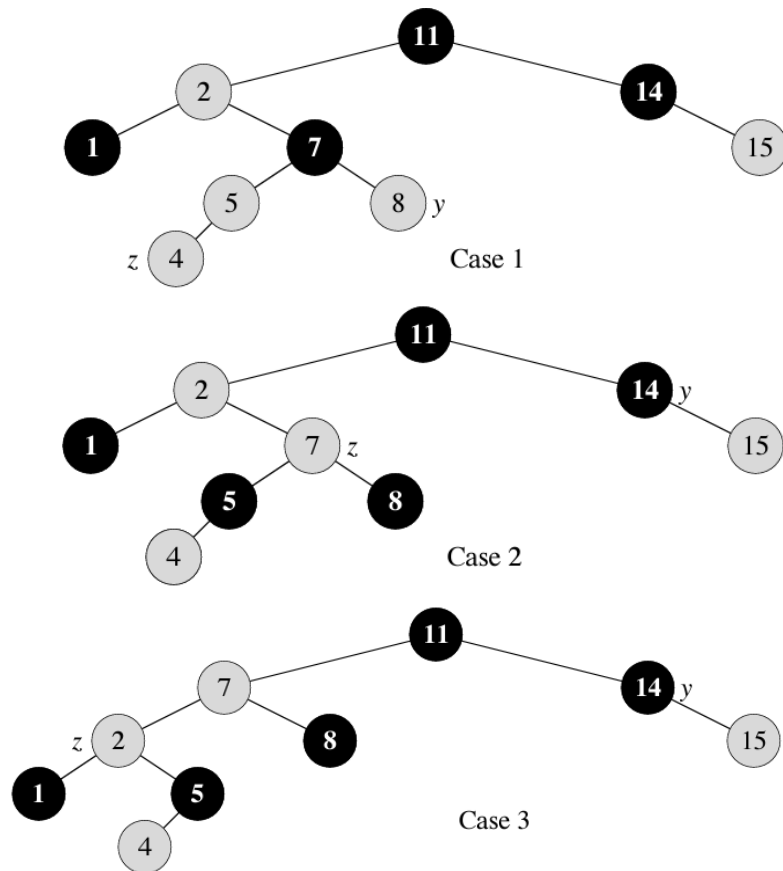
## Actual Code: Finding Spot



RB-INSERT( $T, z$ )

```
1   $y = T.nil$ 
2   $x = T.root$ 
3  while  $x \neq T.nil$ 
4       $y = x$ 
5      if  $z.key < x.key$ 
6           $x = x.left$ 
7      else  $x = x.right$ 
8   $z.p = y$ 
9  if  $y == T.nil$ 
10      $T.root = z$ 
11  elseif  $z.key < y.key$ 
12      $y.left = z$ 
13  else  $y.right = z$ 
14   $z.left = T.nil$ 
15   $z.right = T.nil$ 
16   $z.color = RED$ 
17  RB-INSERT-FIXUP( $T, z$ )
```

# Actual Code: Fix-Up



RB-INSERT-FIXUP( $T, z$ )

```

1  while  $z.p.color == \text{RED}$ 
2      if  $z.p == z.p.p.left$ 
3           $y = z.p.p.right$ 
4          if  $y.color == \text{RED}$ 
5               $z.p.color = \text{BLACK}$ 
6               $y.color = \text{BLACK}$ 
7               $z.p.p.color = \text{RED}$ 
8               $z = z.p.p$ 
9          else if  $z == z.p.right$ 
10              $z = z.p$ 
11             LEFT-ROTATE( $T, z$ )
12              $z.p.color = \text{BLACK}$ 
13              $z.p.p.color = \text{RED}$ 
14             RIGHT-ROTATE( $T, z.p.p$ )
15         else (same as then clause
16             with “right” and “left” exchanged)
17      $T.root.color = \text{BLACK}$ 
    
```

# Time Complexity Analysis

- Insertion (as a red node):  $O(h)$ , obviously
  - Traversing downward along the path to a leaf.
- Fix-up (resolving consecutive reds):  $O(h)$  too!
  - Traversing upward along the path to root at most once.
  - In each iteration of RB-INSERT-FIXUP( $T, z$ )'s while loop,
    - There are fixed number of operations
    - Each iteration pushes up  $z$  one level up
    - The loop can iterate at most all the way up to root, which is  $h$  times.
- Therefore,  $O(h) = O(\lg n)$  all the time (incl. worst case)!

# Deleting Existing Value from a Red-Black Tree

Allow Extra Black On Any Node, Bubble It Up Or Pass It Over

Still Very Complicated!


# Lesson Objectives

- Identify the resulting red-black tree after deleting an arbitrary existing value from a given red-black tree.
- Given an incomplete Red-Black tree deletion code, fill in the blanks for correct Red-Black tree deletion operations.



# Problem Reduction Of RB-DELETE

- Only consider cases of deleting a node with at most one non-NIL child
  - If the value to be deleted is found at a node with two children,
    - Find its successor node (which can't have a left child)
    - Copy the successor value to the original node to be deleted
    - Then delete the successor node (move up its right child to its position)




DELETE( $p$ )  $\rightarrow$  Copy  $p$ 's successor's value to  $p$ 's node, then DELETE( $p$ 's successor node)

Note: In CLRS, it's not copying & deleting, but TRANSPLANTing twice! ( $s$ .right to  $s$ ,  $s$  to  $p$ )

# More Problem Reduction


- Deleting a red node is straightforward
  - In fact, if the node to be deleted is red (with at most one child), then it must be a value-bearing leaf (with no value-bearing child). Can be easily removed without violating any red-black properties
- Deleting a black node is complicated
  - Move up its right child (from previous slide)
  - However, the deleted black node's "black" color has to stay.
    - Giving an **extra black** to the node that's moved up to the deleted black node
    - This extra black needs to be fixed up.



- Note that q might be NIL!
- If q is NIL, then of course q is black, and the NIL now has extra black.
- If q is not NIL and red, then it can be simply recolored to black and we are done.
- Once this step is done, there's no more transplanting and the extra black will need to be fixed.

# How To Fix Extra Black


- Case 1: Extra black node's sibling is red
  - The red sibling must have two black children (same-black-heights rule) and its parent must be black too (no-consecutive-reds rule).
  - Rotate the tree along the parent and make the extra black node's sibling black (Transform to Case 2)



Recolor B & D so that the black-height property is still met.

Note: Any of A, C, and E may be NIL!


- Case 2: Extra black node's sibling and its two children are all black
  - We can recolor the sibling to red, and bubble up the extra black to the parent.
  - If the parent was originally red, it can be simply recolored to black, and we are done. Otherwise, continue fixing up the extra black.



- Black-height requirement is still satisfied
- And the consecutive reds can be easily fixed by recoloring with the extra black
- Or the extra black can be bubbled up/passed along again (iteration)

Note: Any of A, C, and E may be NIL!


- Case 3: Extra black node's sibling and only its right child is black
  - Rotate along the sibling, and recolor so that it's transformed to Case 4.



- Recolor  $C$  &  $D$  so that the black heights are still the same along any path to a leaf.
- This is Case 4.

Note: Any of  $A$ ,  $\gamma$ ,  $\delta$  and  $E$  may be NIL!

- Case 4: Extra black node's sibling is black and its right child is red.
  - Rotate & recolor (maintaining the same black heights)
  - The extra black is gone, and we are done.



Recolor B, D, E as necessary to retain the same black heights and to remove the extra black.

Note: Any of A and C may be NIL!

RB-DELETE( $T, z$ )

```
1   $y = z$ 
2   $y\text{-original-color} = y.\text{color}$ 
3  if  $z.\text{left} == T.\text{nil}$ 
4       $x = z.\text{right}$ 
5      RB-TRANSPLANT( $T, z, z.\text{right}$ )
6  elseif  $z.\text{right} == T.\text{nil}$ 
7       $x = z.\text{left}$ 
8      RB-TRANSPLANT( $T, z, z.\text{left}$ )
9  else  $y = \text{TREE-MINIMUM}(z.\text{right})$ 
10      $y\text{-original-color} = y.\text{color}$ 
11      $x = y.\text{right}$ 
12     if  $y.p == z$ 
13          $x.p = y$ 
14     else RB-TRANSPLANT( $T, y, y.\text{right}$ )
15          $y.\text{right} = z.\text{right}$ 
16          $y.\text{right}.p = y$ 
17     RB-TRANSPLANT( $T, z, y$ )
18      $y.\text{left} = z.\text{left}$ 
19      $y.\text{left}.p = y$ 
20      $y.\text{color} = z.\text{color}$ 
21 if  $y\text{-original-color} == \text{BLACK}$ 
22     RB-DELETE-FIXUP( $T, x$ )
```

Actual Code:

RB-TRANSPLANT(), RB-DELETE()

RB-TRANSPLANT( $T, u, v$ )

```
1  if  $u.p == T.\text{nil}$ 
2       $T.\text{root} = v$ 
3  elseif  $u == u.p.\text{left}$ 
4       $u.p.\text{left} = v$ 
5  else  $u.p.\text{right} = v$ 
6       $v.p = u.p$ 
```

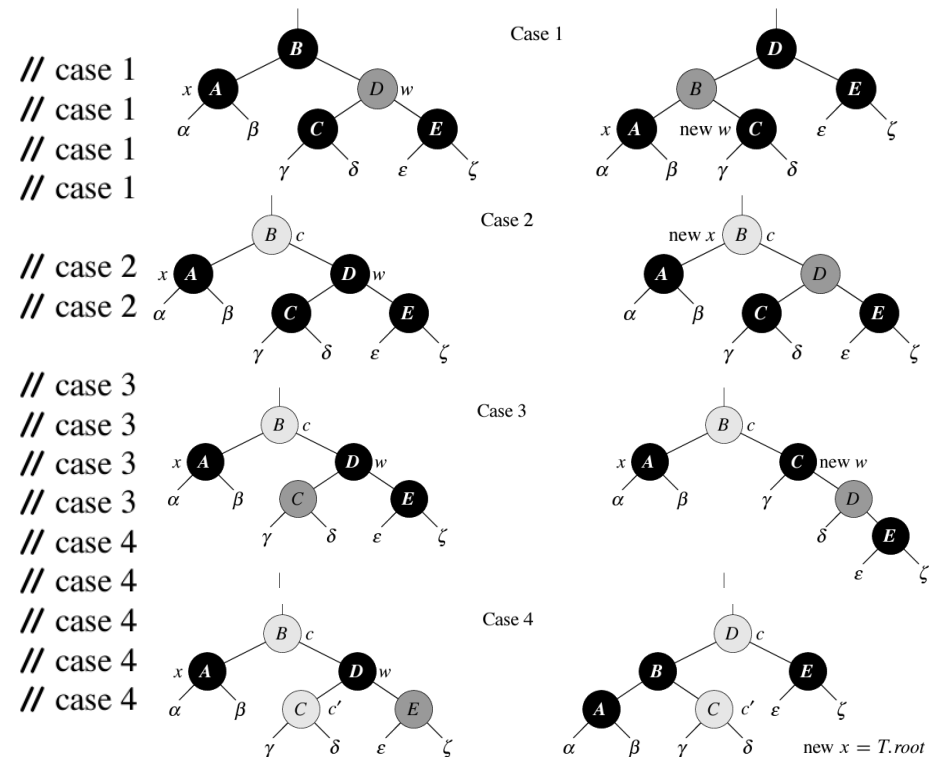
## RB-DELETE-FIXUP( $T, x$ )

```

1  while  $x \neq T.root$  and  $x.color == BLACK$ 
2      if  $x == x.p.left$ 
3           $w = x.p.right$ 
4          if  $w.color == RED$ 
5               $w.color = BLACK$ 
6               $x.p.color = RED$ 
7              LEFT-ROTATE( $T, x.p$ )
8               $w = x.p.right$ 
9          if  $w.left.color == BLACK$  and  $w.right.color == BLACK$ 
10              $w.color = RED$ 
11              $x = x.p$ 
12         else if  $w.right.color == BLACK$ 
13              $w.left.color = BLACK$ 
14              $w.color = RED$ 
15             RIGHT-ROTATE( $T, w$ )
16              $w = x.p.right$ 
17              $w.color = x.p.color$ 
18              $x.p.color = BLACK$ 
19              $w.right.color = BLACK$ 
20             LEFT-ROTATE( $T, x.p$ )
21              $x = T.root$ 
22         else (same as then clause with “right” and “left” exchanged)
23      $x.color = BLACK$ 

```

## RB-DELETE-FIXUP()





# Time Complexity Analysis

- Still  $O(h) = O(\lg n)$ .
  - Case 3 & 4: Fixed # operations & terminates
  - Case 1: Fixed # operations, transforms to Case 2.
  - Case 2: Fixed # operations,
    - Then terminates if the extra black node's parent is red
    - Or else repeat, but one-level up, meaning it can repeat only up to  $h$  many times.
  - Thus  $O(h)$ !
- We achieved  $O(\lg n)$  for all operations in all cases (incl. worst)

## There are Other Balanced Trees

- AVL Trees: BSTs balanced by height
  - For any node, subtrees should have height difference of at most 1
  - This implies that the worst tree has size  $n(h) = n(h-1) + n(h-2) + 1$
  - Then  $h \leq 1.44 \log_2(n)$
- Weight-balanced BSTs
  - For any node, subtrees should have a bounded size ratio:  
$$1 - \beta \leq \text{size}(\text{left})/\text{size}(\text{right}) \leq 1 + \beta$$
  - Then  $h \leq c(\beta) \log_2(n)$