

CS5800: B⁺-Trees

Ricardo Baeza-Yates & Anurag Bhardwaj

Northeastern University 2018

Agenda

- B⁺-Trees
 - Introduction
 - Searching, Insertion, Deletion
 - Complexity Analysis
- B⁺-Tree Advantages

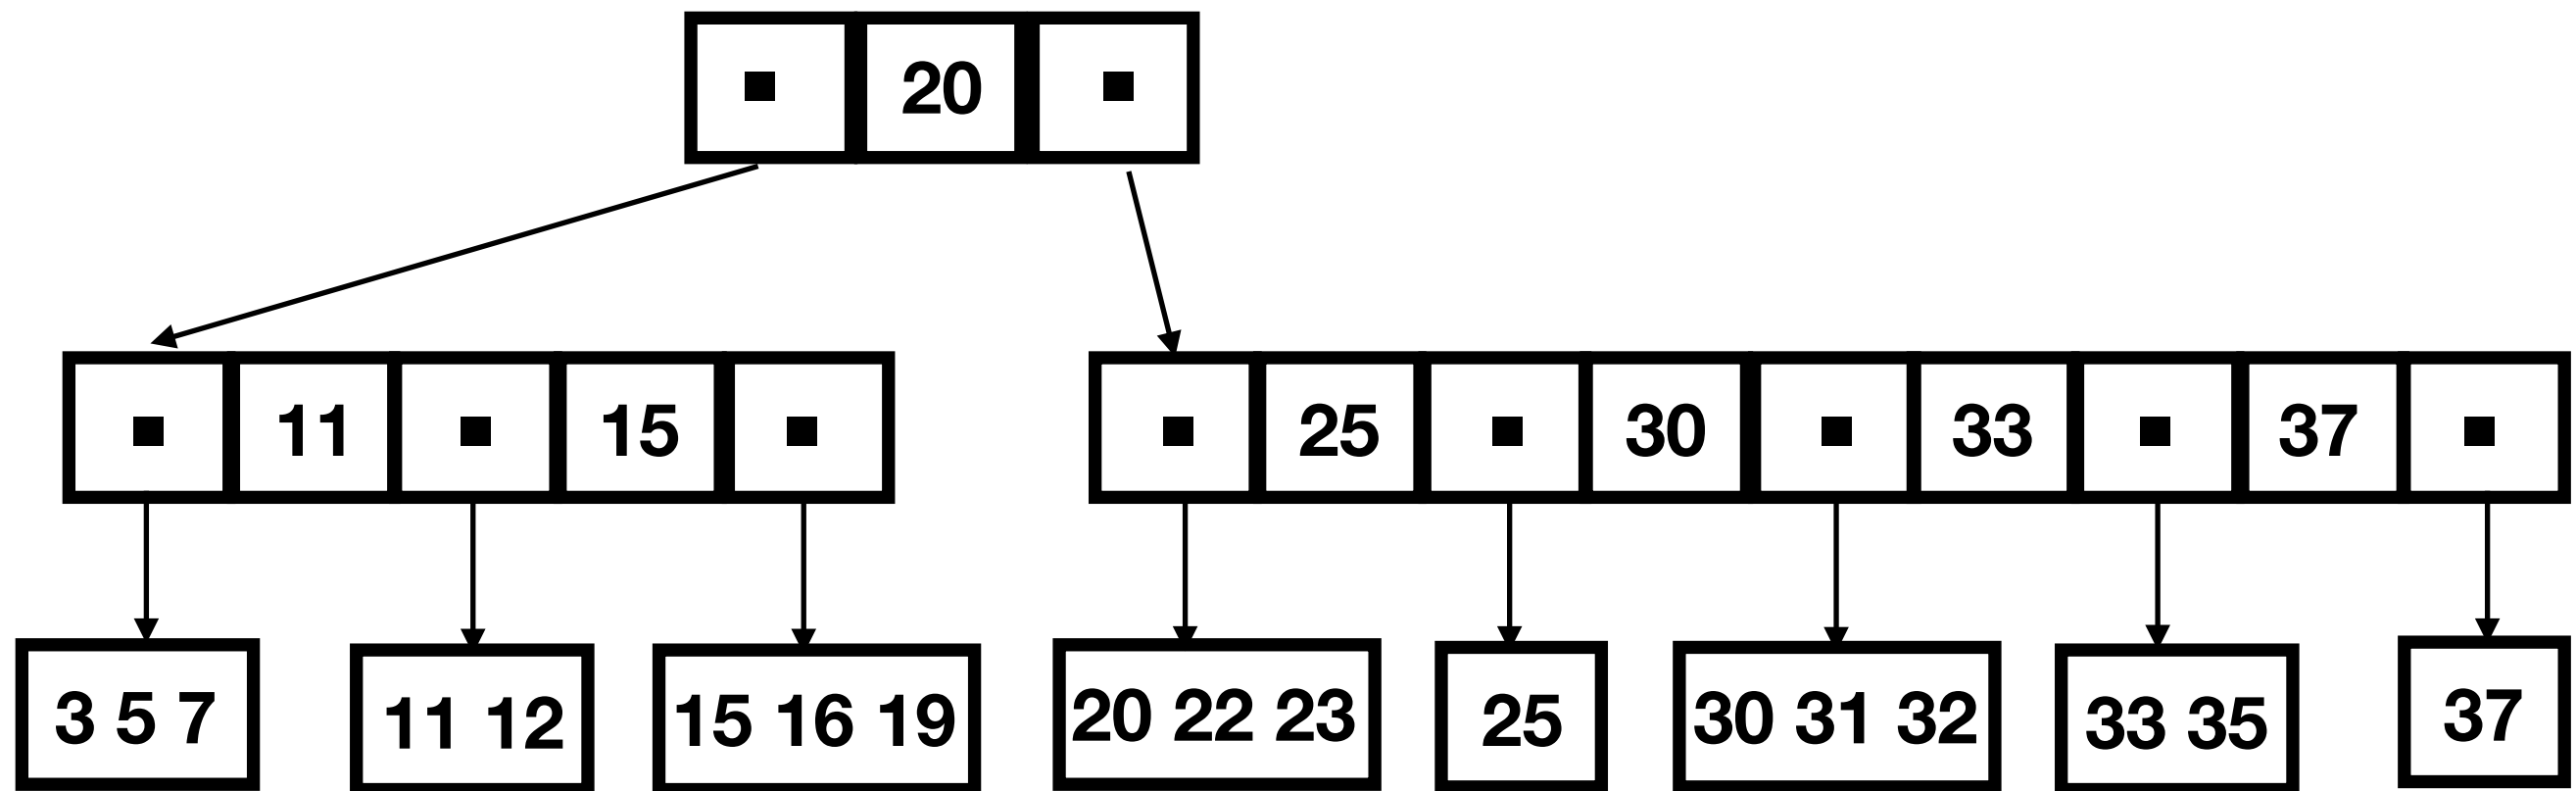
B⁺-Trees - Introduction

- Disk Storage Problem
 - How to store data on disk for faster access?
 - How about Binary Search trees?
 - # of reads needed?
 - # of root -> leaf traversals?
 - Locality of reference?
- B-Trees and its variants comes to rescue
- B-Tree is a variant of search tree with many children

B⁺-Trees - Properties

- Key Ideas
 - Reduce # of reads: store more data in each node
 - Reduce # of traversals: high branching factor ($B \gg 2$)
 - Locality of reference: sorted keys

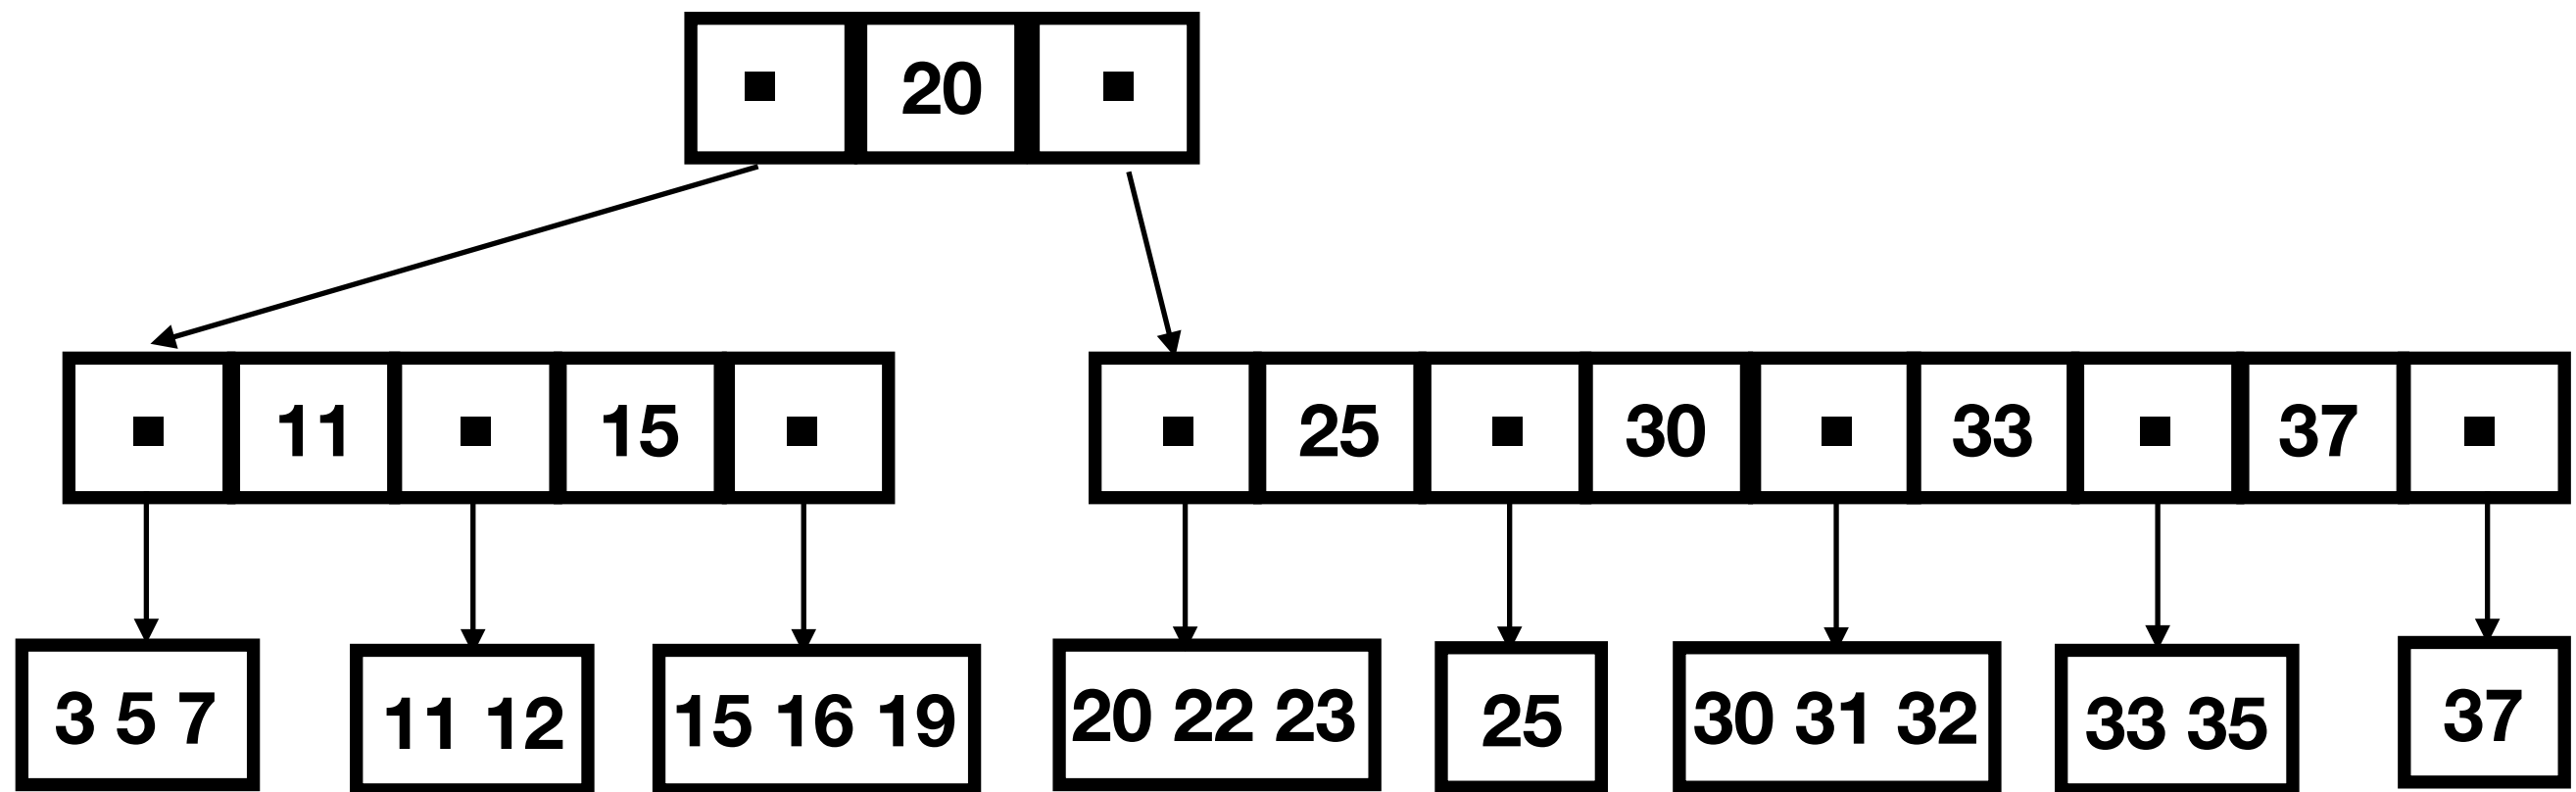
B+-Trees - Example



B⁺-Trees - Formal Definition



- For non-leaf node N:
 - $(\# \text{ of children in } K) - (\# \text{ of keys stored in } K) == 1$
- For every non-root node N:
 - at least $(B-1)$ keys
- All nodes have at most $2*B$ children
- All leaf nodes have same depth
- Only leaf node(s) contain data, non-leaf contain routing values

B⁺-Trees - Search





- Linear or binary search in each node - given max. # of keys in node
- If not found, follow the appropriate child pointer
- Complexity: $\log_B(N)$

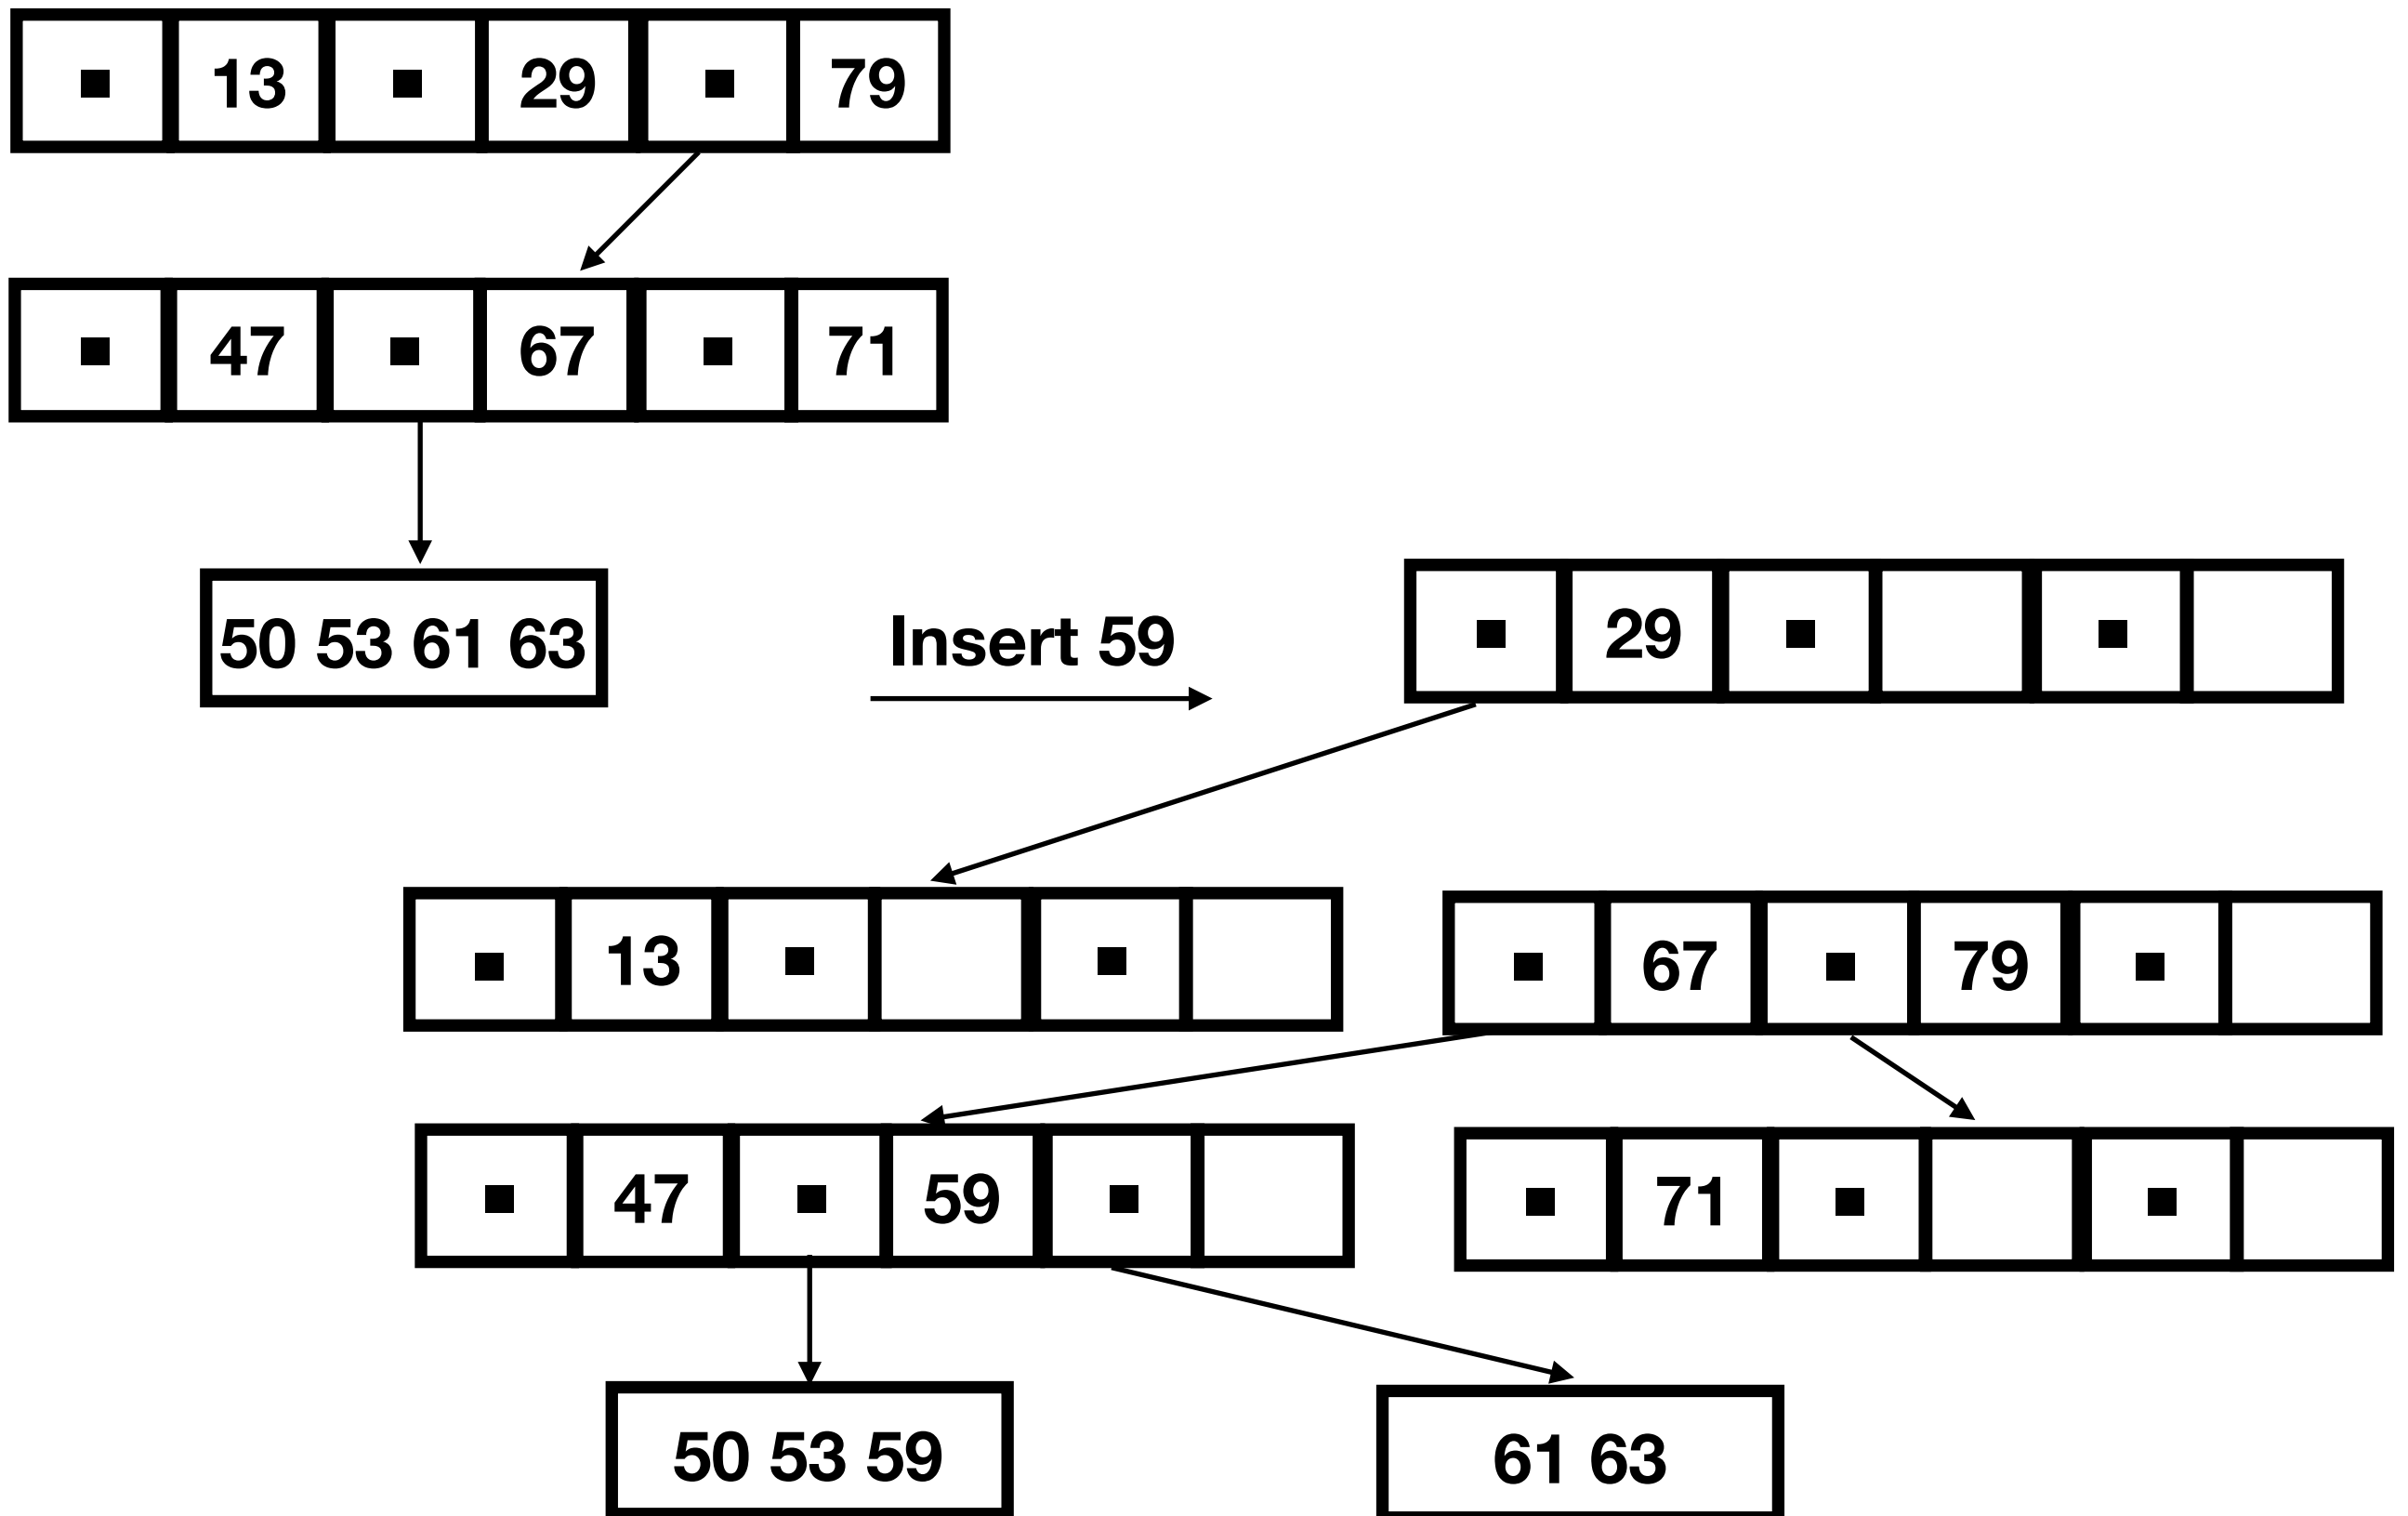
B+-Trees - Insertion (1/3)



B+-Trees - Insertion (2/3)



B+-Trees - Insertion (3/3)



B⁺-Trees - Deletion

- An element is removed from the leaf
 - if leaf becomes empty, remove key from parent
 - Keep adjusting parent and its siblings to satisfy invariants
 - May ripple to root and reduce height of tree by 1
 - If the node has $< 50\%$ occupancy, merge it with left/right node (standard B⁺-Trees have a minimal occupancy of 50%)
- More examples:
 - <https://www.cs.usfca.edu/~galles/visualization/BTree.html>

B⁺-Trees - Complexity

- Time Complexity
 - Search: $O(B * \log_B(n))$
 - Insertion: $O(B * \log_B(n))$
 - Deletion: $O(B * \log_B(n))$
- Space Complexity
 - $O(n)$

B⁺-Trees - Advantages

- Multi-level Indexing
- Bulk Insertion
- Aggregation queries
- Value of B can be set to match the size of a disk sector or its multiples