

CS5800: External Hashing

Ricardo Baeza-Yates & Anurag Bhardwaj

Northeastern University 2018

Agenda

- External Hashing
 - Introduction
 - Collision Avoidance
 - Bucketed Hashing
 - Static vs Dynamic Hashing
- Dynamic Hashing
- Linear Hashing
- Extendible Hashing

Introduction

- External Hashing: Hashing for disk files
- Key Ideas:
 - Target Address Space consists of buckets
 - Each bucket is either a single or multiple contiguous blocks
 - Hash function maps a key to a bucket number which points to an absolute block address on disk
- External Hashing Vs. B+-Trees:
 - Faster than B+-Trees for searches
 - Cannot support range queries

Collision Avoidance

- Collision is not a big problem for external hashing
 - Each bucket can hold as many records as possible
 - When bucket is full, chaining can be used
 - Record pointers chain the bucket to an Overflow Bucket
 - Hashing performance depends on number of overflow buckets

Bucketed Hashing

- Hash Function H generates a bucket address
- Store more than one key at the bucket address
- N Hash addresses split into B buckets
- Each bucket gets N/B slots
- Overflow record handling

Bucketed Hashing

- Insertion:
 - Given key k , generated bucket address $b = h(k)$
 - Find the first empty slot at b from $(1, 2, \dots, N/B)$ and store key there
 - If all N/B slots are full, store key in an overflow bucket
- Search:
 - Given a search key sk , generate bucket address $b = h(sk)$
 - Search for key sk at bucket b , return if found, return NF if free slots
 - If sk is not found at b , search in overflow bucket, return if found
 - If sk is not found in overflow bucket, return NF

Bucketed Hashing

- Deletion:
 - An important step to ensure correct search behavior
 - e.g. Insert 4 keys in sequence {A, J, M, S} to a given bucket b
 - Hash Assignment: 1 -> A, 2 -> J, 3 -> M, 4 -> S
 - Delete M, new assignment: 1 -> A, 2 -> J, 3 -> empty, 4 -> S
 - Search S, will return empty since search terminated at 3
 - Solution: mark 3 as TOMBSTONE and continue search
 - Correct assignment: 1 -> A, 2 -> J, 3 -> TOMBSTONE, 4 -> S
 - Modify insertion to add record at TOMBSTONE as well


Static vs. Dynamic Hashing

- Static Hashing
 - Number of buckets is fixed
 - Expensive to re-organize or re-hash the whole file
- Dynamic Hashing
 - Number of buckets grows/shrinks dynamically
 - Low cost

Dynamic Hashing (1/2)

- Core Idea
 - Start with 1 bucket to store the records
 - Continue till bucket is full
 - If new record is added, split 1 bucket into 2 buckets
 - Records are distributed based on 1-MSB of binary hash value
- Directory
 - Internal node to guide the search
 - Leaf node to point to bucket

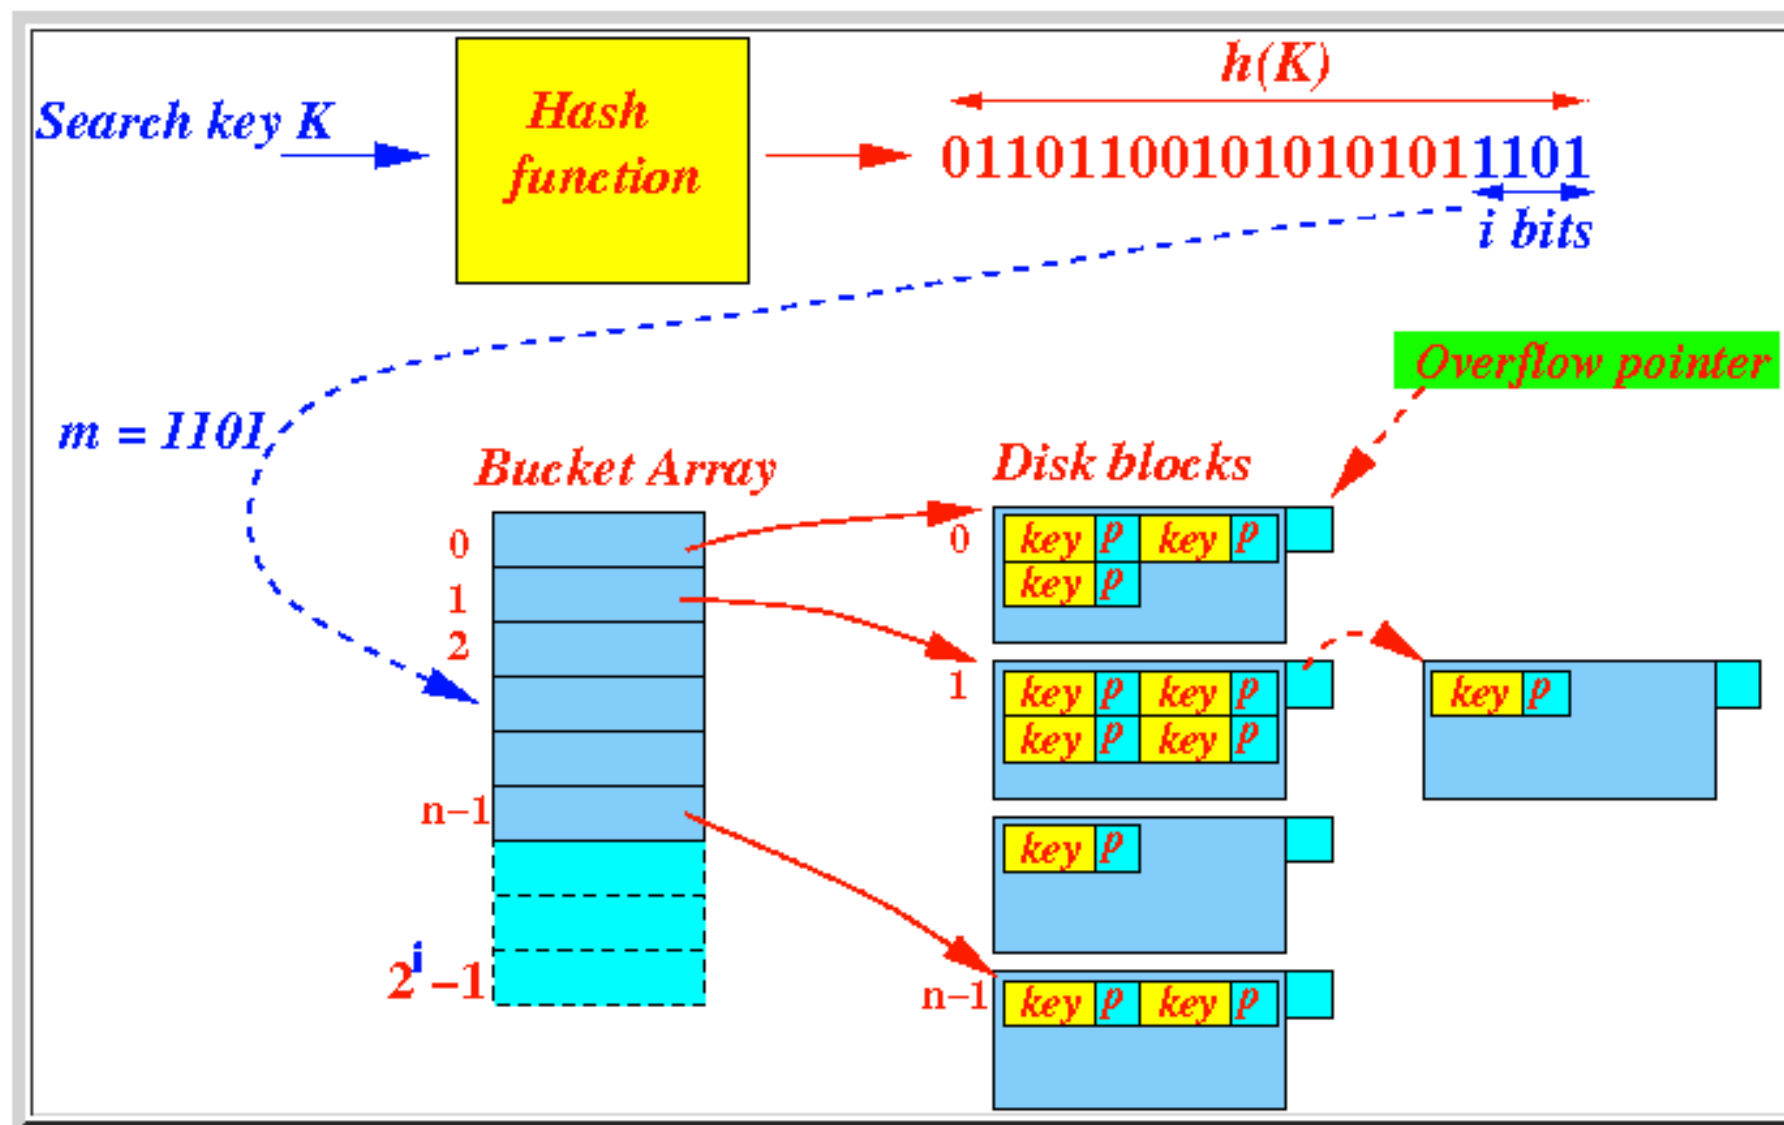
Dynamic Hashing (2/2)



Linear Hashing

- Core Idea
 - Allows dynamic shrinking/growing of buckets without directory
 - When avg. occupancy per bucket $>$ threshold, split happens
 - Growth rate of bucket is linear
 - Initial hash function: $h_0 = h(k) = k \bmod M$
 - Uses a family of hash functions
 - Example shown next

Linear Hashing



- Core Idea

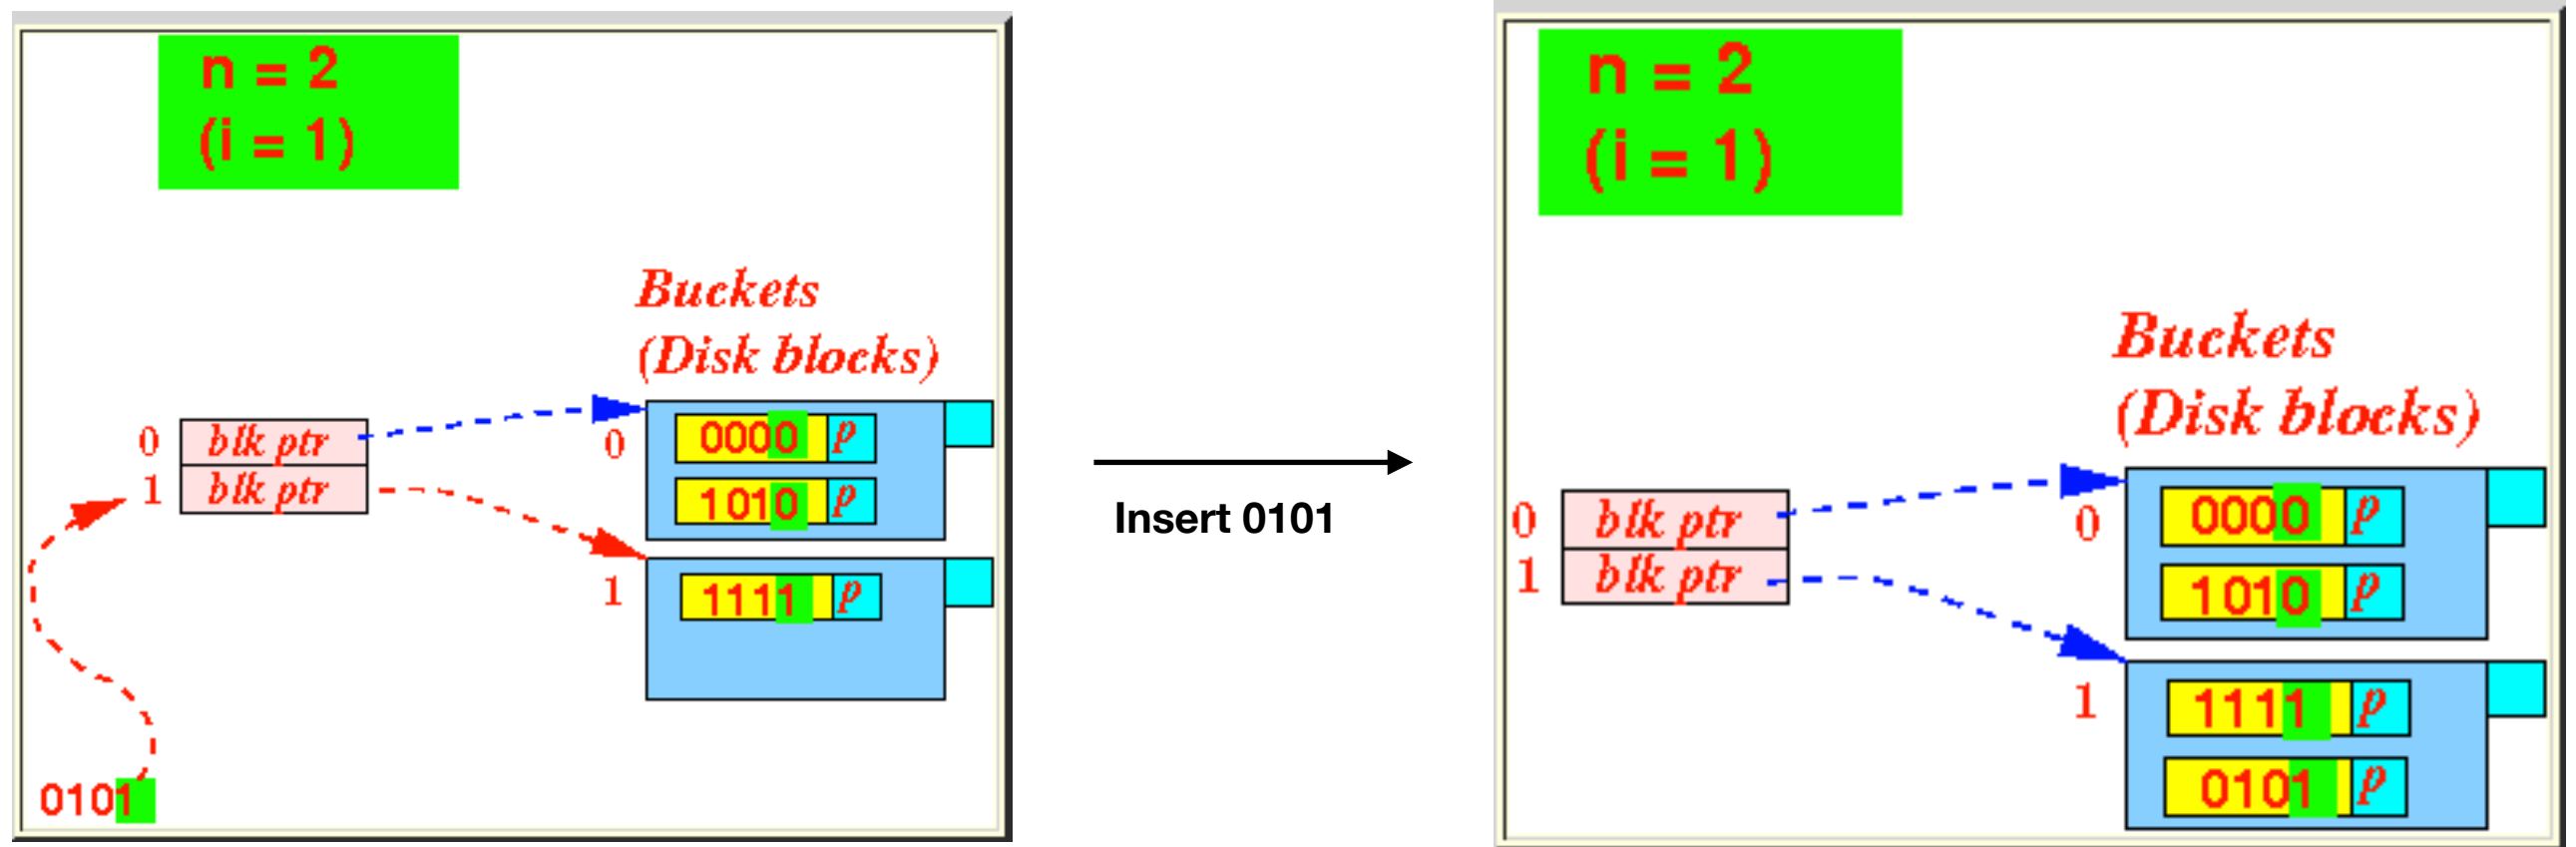
- “ n ” real buckets, last “ i ” bits from hash address are used
- “ i ” bits can address 2^i buckets which can be $> n$
- Bucket numbers $> (n-1)$ are called virtual buckets and have their first bit as 1, 1XXXXXX

<http://www.mathcs.emory.edu/~cheung/Courses/554/Syllabus/3-index/linear-hashing.html>

Linear Hashing

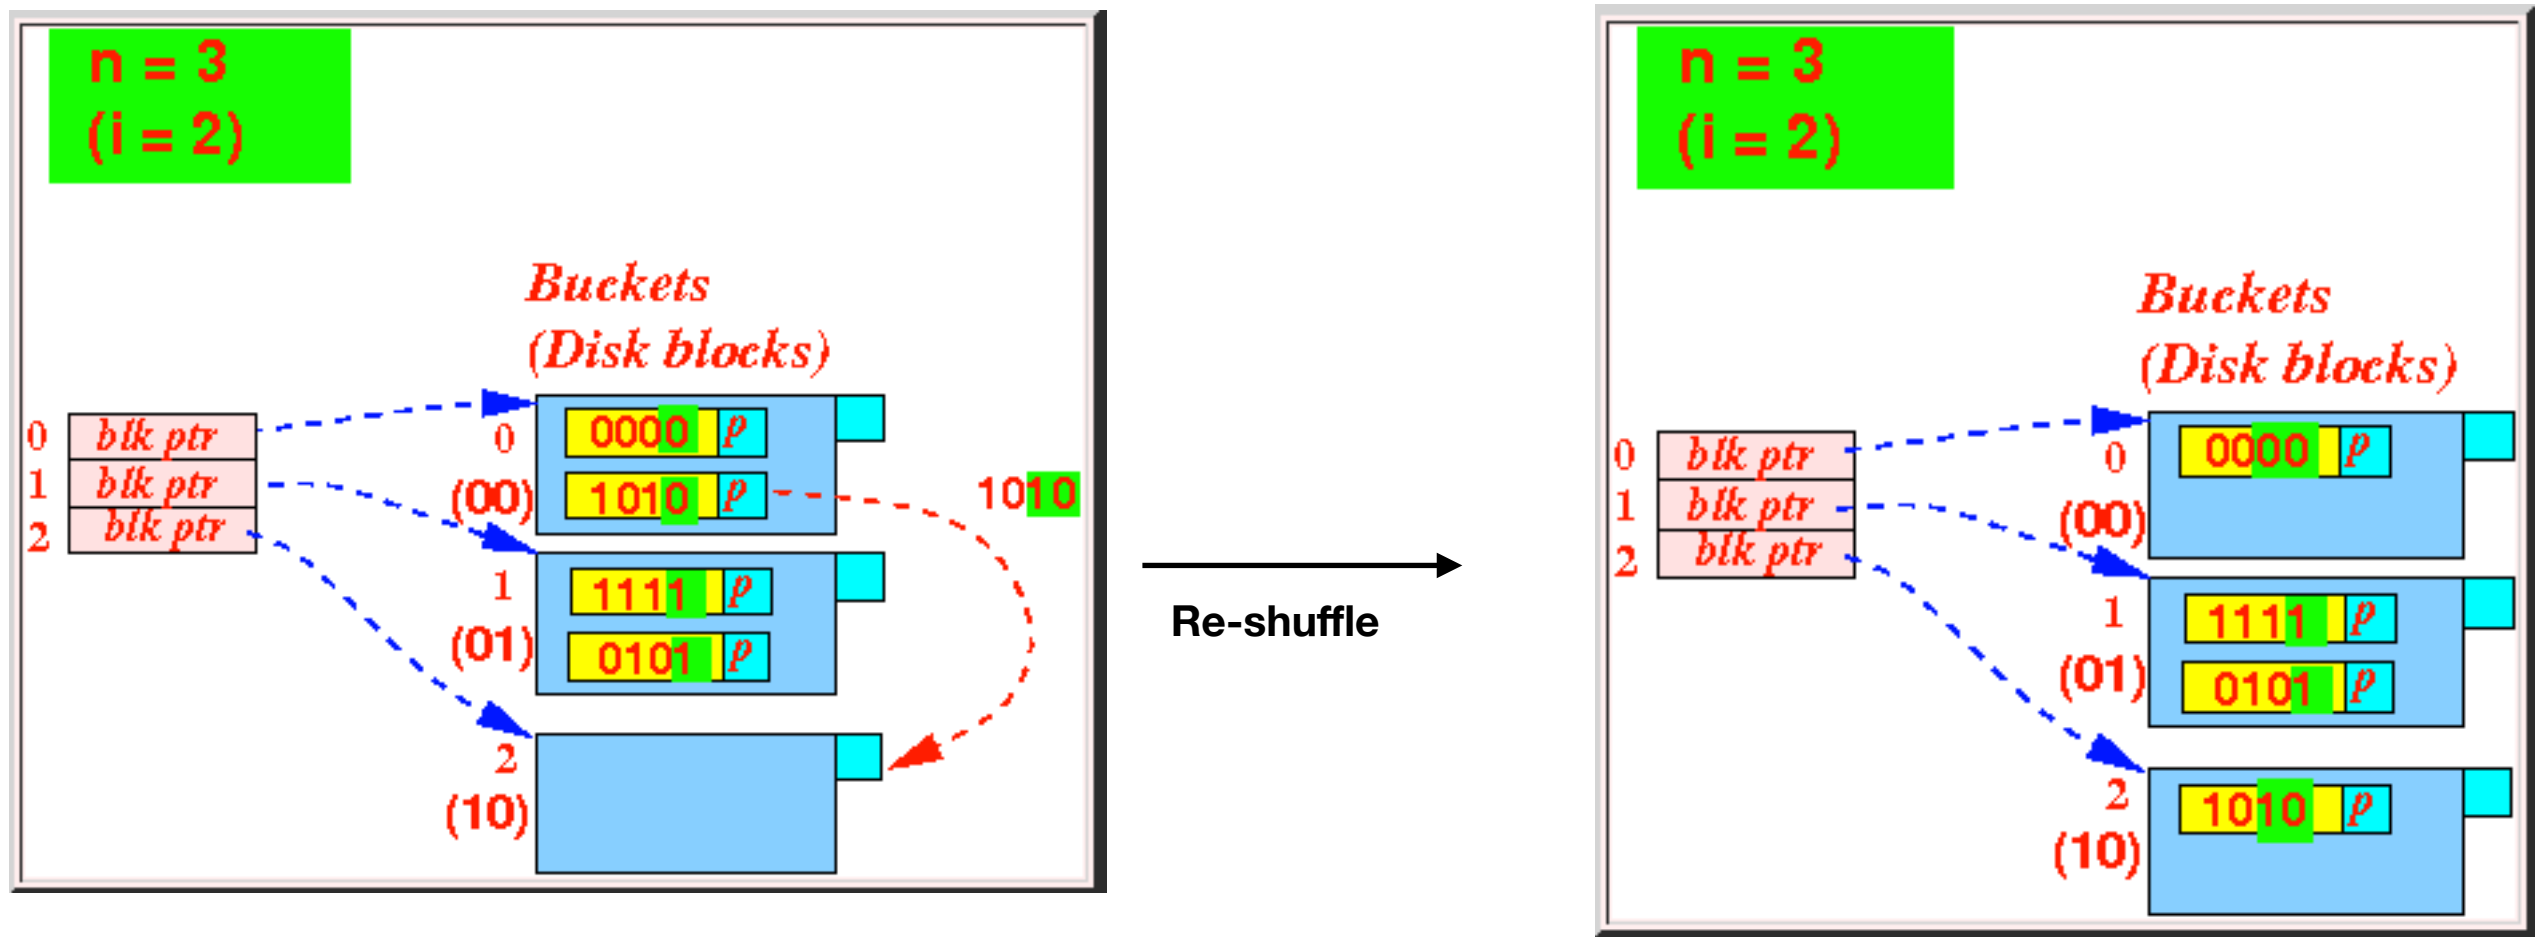
- Split Criteria
 - r is the number of records
 - n is the number of buckets
 - b is the bucket size or number of keys that can be stored in bucket
 - Avg occupancy = $\frac{r}{n * b}$
 - Split when avg. occupancy > some threshold
 - Increase criteria: $n = n + 1$

Linear Hashing Example



- Avg Occupancy before Insertion = $3 / (2 \times 2) = 0.75$
- Avg Occupancy after insertion = $4 / (2 \times 2) = 1$
- Let's assume split threshold = 0.85
- Since $1 > 0.85$, we need to split and add a block

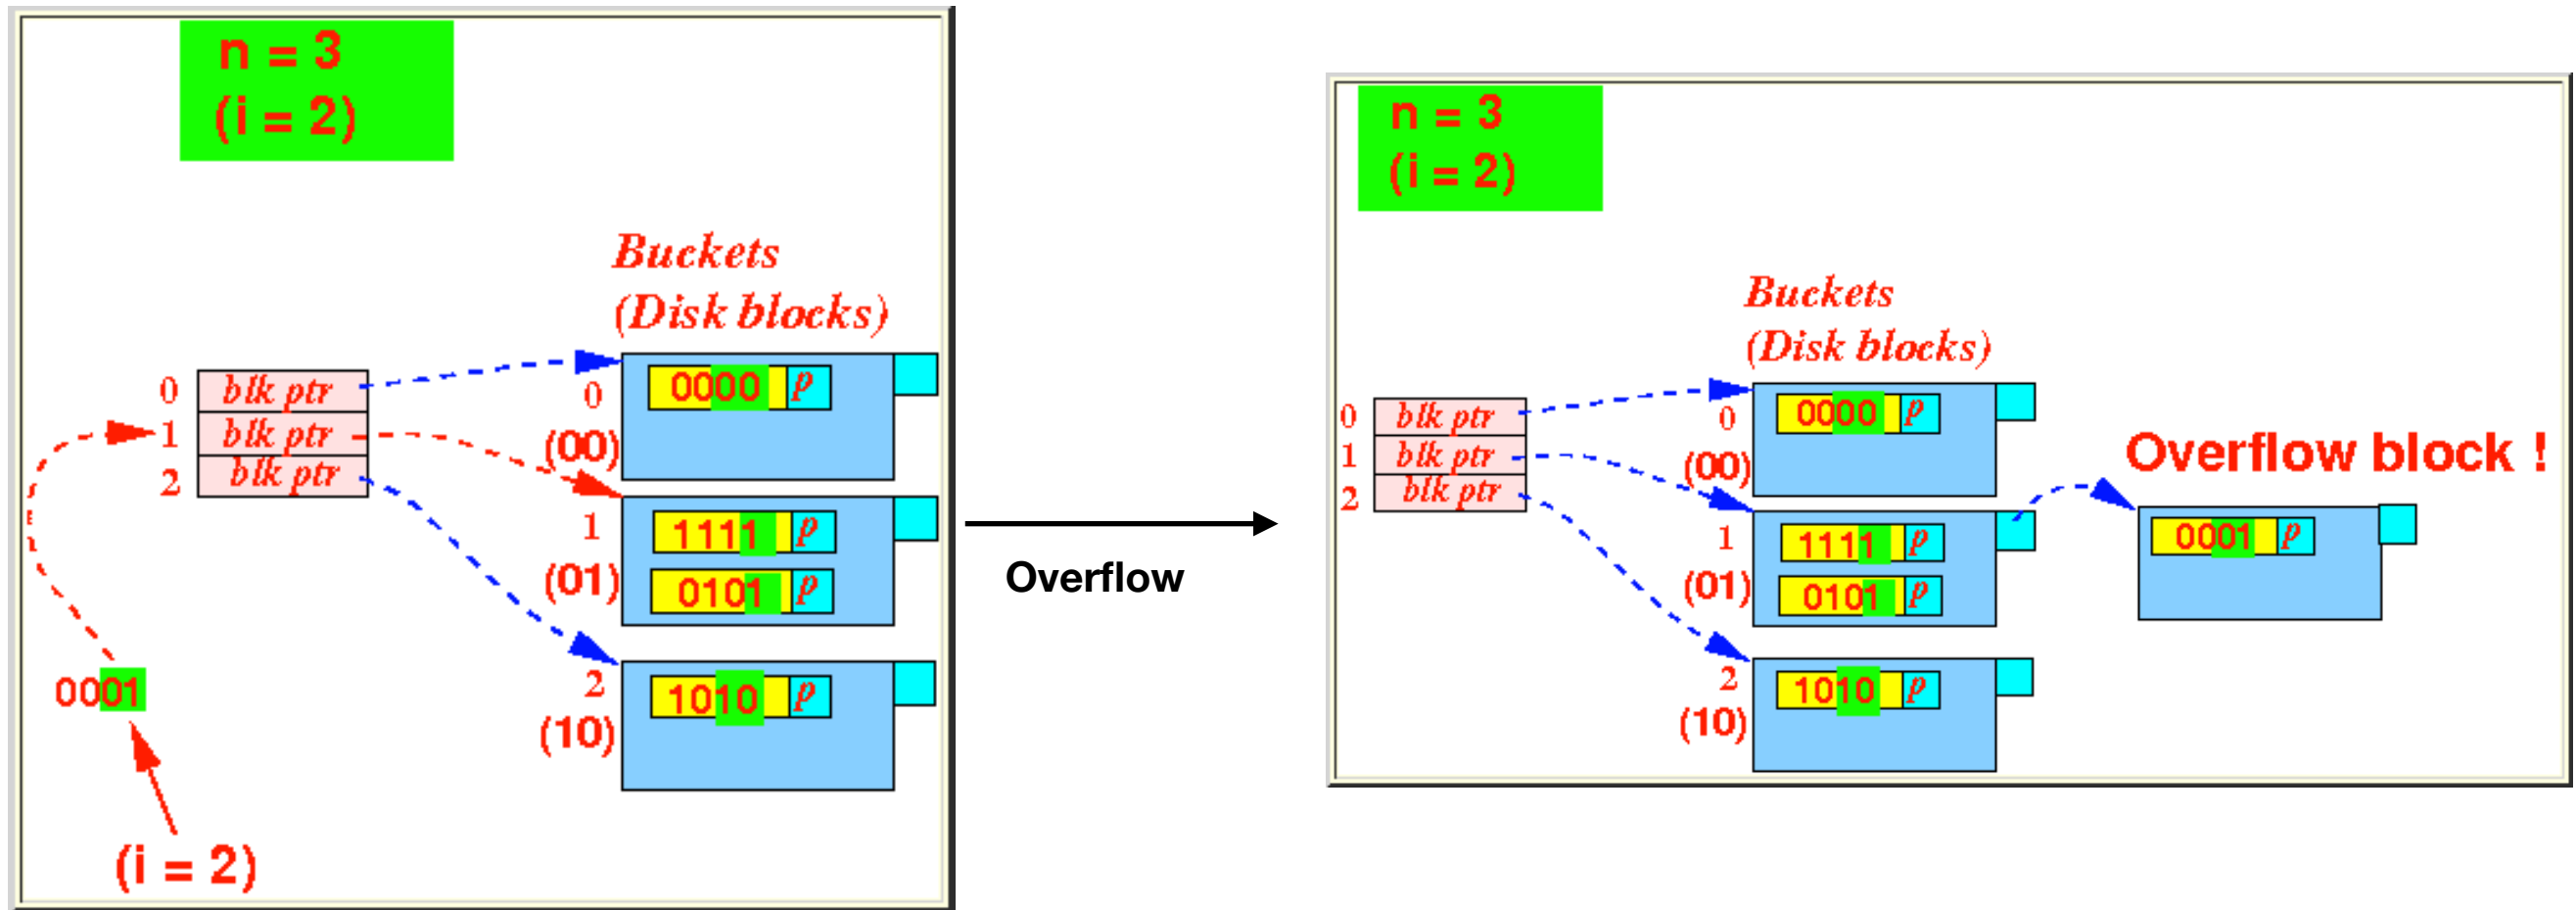
Linear Hashing Example



- Add bucket 2 with id 01, notice i changes from 1 to 2 since we now need 2 bits to address buckets
- Transfer key from bucket 00 to bucket 10
- Avg. occupancy after the split and shuffle: $4 / (3 \cdot 2) = 0.67 < 0.85$
- How to search key 1111 ?
 - $i = 2$, bucket address 11 does not exist or is a virtual bucket, Flip the MSB from 1 to 0, 11 \rightarrow 01

<http://www.mathcs.emory.edu/~cheung/Courses/554/Syllabus/3-index/linear-hashing.html>

Linear Hashing Example

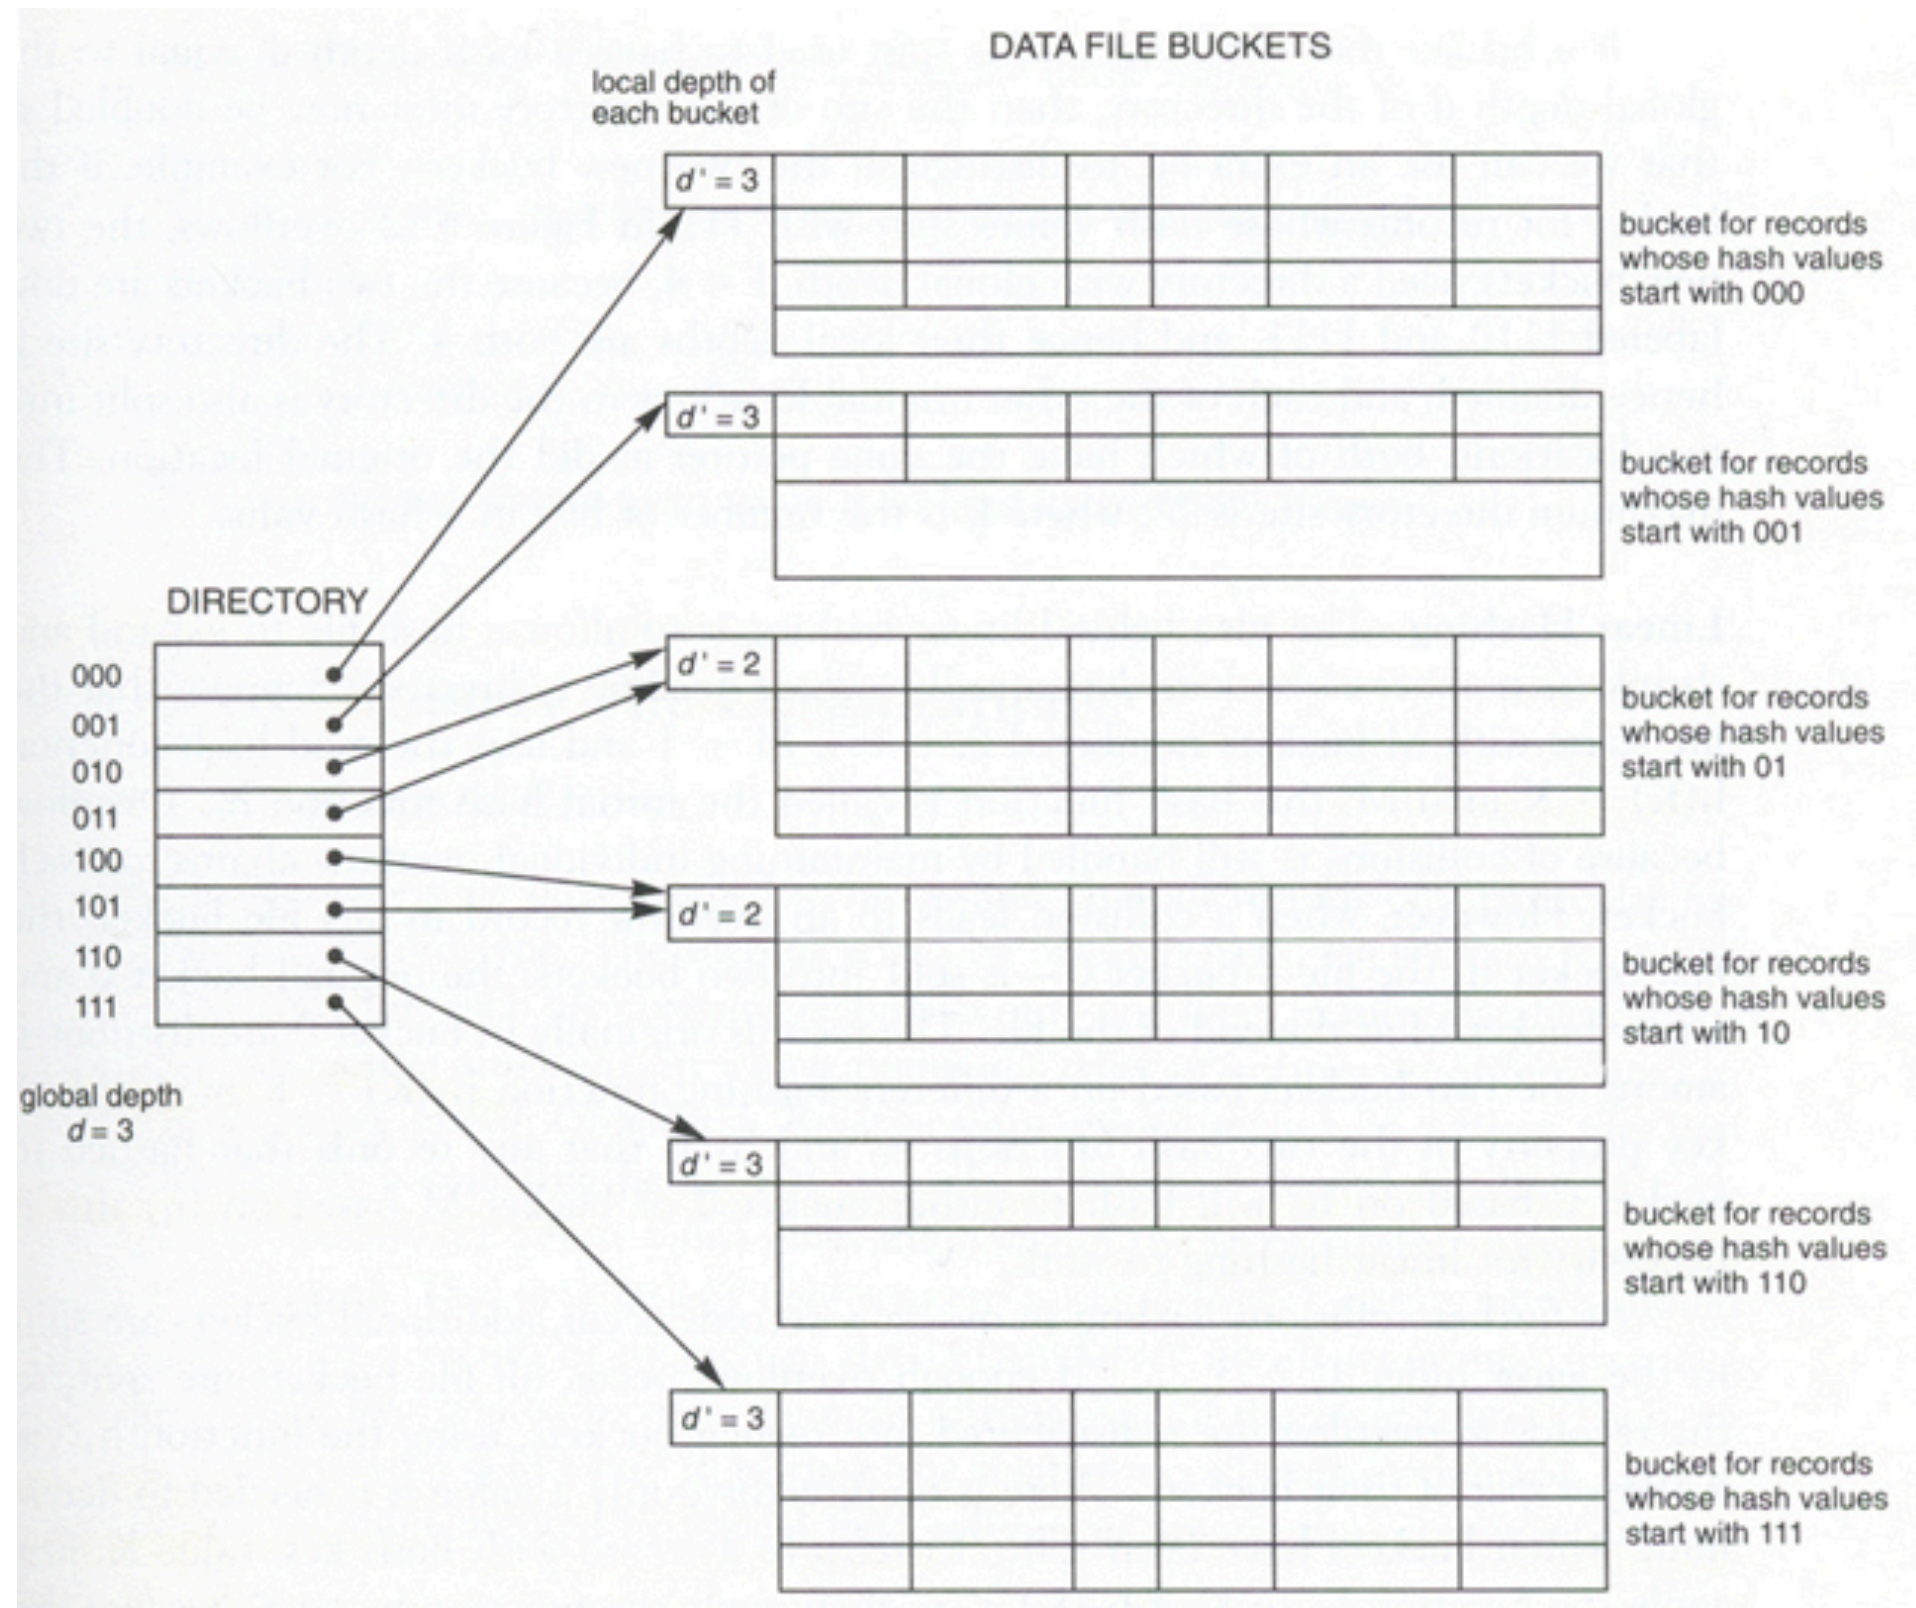


- Avg. occupancy after overflow: $5 / (3 \cdot 2) = 0.83 < 0.85$

Extendible Hashing (1/2)

- Core Idea
 - Variant of dynamic hashing with different directory structure
 - Directory contains 2^d addresses; d = Global depth
 - Local depth = d'
 - Local depth determines number of bits to use for hash values
 - If a bucket fills up, a split happens
 - Example shown next
 - Issue: exponential increase in size of hash table

Extendible Hashing (2/2)



**When this bucket fills up,
it splits into 2 buckets of local depth = 3
containing hash values 010 and 011**